

# Worksheet-5

```
[34] In [1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
from sklearn.model_selection import train_test_split  
  
[34] In [2]: student_df = pd.read_csv('/content/drive/MyDrive/Concepts and Technologies of AI/student_week_5.csv')  
  
[34] In [3]: print("First five rows of the data\n", student_df.head())  
  
[34] Out[3]:  
First five rows of the data  
   Math  Reading  Writing  
0    48       68      63  
1    62       81      72  
2    73       80      70  
3    76       83      79  
4    59       64      62
```

```
[34] In [4]: print("Last five rows of the data\n", student_df.tail())  
  
[34] Out[4]:  
Last five rows of the data  
   Math  Reading  Writing  
995    72       74      78  
996    73       86      90  
997    89       87      94  
998    81       82      78  
999    66       66      72  
  
[35] In [5]: print("Description of the data\n", student_df.describe())  
  
[35] Out[5]:  
Description of the data  
          Math      Reading      Writing  
count  1000.000000  1000.000000  1000.000000  
mean   67.290000  80.722000  68.160000  
std    15.485000  14.657027  15.241287  
min    13.000000  19.000000  14.000000  
25%   58.000000  69.750000  58.000000  
50%   68.000000  70.000000  69.500000  
75%   78.000000  81.000000  79.000000  
max   100.000000  107.000000  100.000000  
  
[36] In [6]: X = student_df[['Math', 'Reading']].values #features  
Y = student_df['Writing'].values #target
```

```
[36] In [7]: print("X (Feature matrix) shape:", X.shape)  
print("Y (Target vector) shape:", Y.shape)  
  
[36] Out[7]:  
X (Feature matrix) shape: (1000, 2)  
Y (Target vector) shape: (1000,)  
  
[46] In [8]: X_transposed = X.T  
print("X transposed shape (dnn format):", X_transposed.shape)  
  
[46] Out[8]:  
X transposed shape (dnn format): (2, 1000)
```

```
[41] In [9]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)  
print("Training set size: ", X_train.shape[0], "samples")  
print("Test set size: ", X_test.shape[0], "samples")  
  
[41] Out[9]:  
Training set size: 800 samples  
Test set size: 200 samples
```

**Task - 4**

```
[42] 0 def cost_function(X, Y, W):
    # Parameters:
    # This function finds the Mean Square Error.
    # Input parameters:
    # X: Feature Matrix
    # Y: Target Matrix
    # W: weight matrix
    # Output Parameters:
    # cost: accumulated mean square error.

    n = len(Y)
    Y_pred = np.dot(X, W)
    squared_loss = (Y_pred - Y)**2
    cost = (1 / (2 * n)) * np.sum(squared_loss)
    return cost
```

**Task-5**

```
[43] 0 X_test_case = np.array([[1, 2], [3, 4], [5, 6]])
Y_test_case = np.array([3, 7, 11])
W_test_case = np.array([1, 1])

cost = cost_function(X_test_case, Y_test_case, W_test_case)
print("Cost function output: {cost}")

if cost == 0:
    print("Cost function passed! Proceed further.")
else:
    print("Something went wrong! Reimplement cost function.")

Cost function output: 0.0
Cost function passed! Proceed further.
```

**Task-6**

```
[44] 0 def gradient_descent(X, Y, W, alpha, iterations):
    """
    Perform gradient descent to optimize the parameters of a linear regression model.

    Parameters:
    # X (numpy.ndarray): Feature matrix (m x n).
    # Y (numpy.ndarray): Target vector (m x 1).
    # W (numpy.ndarray): Initial guess for parameters (n x 1).
    # alpha (float): Learning rate.
    # iterations (int): Number of iterations for gradient descent.
    # Returns:
    # tuple: A tuple containing the final optimized parameters (W_update) and the history of cost values
    """
    # Initialize cost history
    n = len(Y)
    cost_history = [0] * iterations
    # Number of samples

    for iteration in range(iterations):
        # Step 1: Hypothesis Values
        Y_pred = np.dot(X, W)
        # Step 2: Difference between Hypothesis and Actual Y
        loss = Y_pred - Y
        # Step 3: Gradient Calculation
        dw = (1/n)*np.dot(X.T, loss)
        # Step 4: Updating Values of W using Gradient
        W = W - alpha * dw
        # Step 5: New Cost Value
        cost = cost_function(X, Y, W)
        cost_history.append(cost)
    return W, cost_history
```

**Task-7**

```
[45] 0
# Generate random test data
np.random.seed(0) # For reproducibility
X = np.random.rand(100, 3) # 100 samples, 3 features
Y = np.random.rand(100)
W = np.random.rand(3) # Initial guess for parameters
# Set learning parameters
alpha = 0.01
iterations = 1000
# Test the gradient_descent function
final_params, cost_history = gradient_descent(X, Y, W, alpha, iterations)
# Print the final parameters and cost history
print("Final Parameters:", final_params)
print("Cost History:", cost_history[:10])

Final Parameters: [0.20551667 0.54295881 0.10388027]
Cost History: [0.  0.  0.  0.  0.  0.  0.  0.  0.]
```

Worksheet-5.ipynb

```
[46] # Model Evaluation - RMSE
def rmse(Y, Y_pred):
    """
    This Function calculates the Root Mean Squares.
    # Input Arguments:
    # Y: Array of actual(Target) Dependent Variables.
    # Y_pred: Array of predicted Dependent Variables.
    # Output Arguments:
    # rmse: Root Mean Square.
    """
    rmse = np.mean((Y - Y_pred) ** 2)
    return rmse
```

Worksheet-5.ipynb

```
[47] # Model Evaluation - R2
def r2(Y, Y_pred):
    """
    This Function calculates the R Squared Error.
    # Input Arguments:
    # Y: Array of actual(Target) Dependent Variables,
    # Y_pred: Array of predicted Dependent Variables.
    # Output Arguments:
    # rsquared: R Squared Error.
    """
    mean_y = np.mean(Y)
    ss_tot = np.sum((Y - np.mean(Y)) ** 2)
    ss_res = np.sum((Y - Y_pred) ** 2)
    r2 = 1 - (ss_res / ss_tot)
    return r2
```

Worksheet-5.ipynb

```
[54] def main():
    # Step 1: Load the dataset
    # Step 2: Split the data into features (X) and target (Y)
    # Step 3: Split the data into training and test sets (80% train, 20% test)
    # Step 4 and 5 are already done above.
    # Step 6: Initialize weights to zeros, learning rate and number of iterations
    W = np.zeros(X_train.shape[1]) # Initialize weights
    alpha = 0.00001 # Learning rate
    iterations = 1000 # Number of iterations for gradient descent
    print("Initial weights: (W)")
    print("Learning rate: (alpha)")
    print("Iterations: (iterations)")
    # Step 5: Perform Gradient Descent
    print()
    M, N = X_train.shape
    cost_history = gradient_descent(X_train, Y_train, W, alpha, iterations)
    # Step 6: Make predictions on the test set
    Y_pred = np.dot(X_test, W_optimal)
    # Step 7: Evaluate the model using RMSE and R-Squared
    model_rmse = rmse(Y_test, Y_pred)
    model_r2 = r2(Y_test, Y_pred)
    # Step 8: Output the results
    print("Final weights: ", W_optimal)
    print("Cost History: (first 10 iterations):", cost_history[:10])
    print("RMSE on Test Set: ", model_rmse)
    print("R-Squared on Test Set: ", model_r2)

    # Plot cost history
    plt.figure(figsize=(10, 6))
    plt.plot(cost_history)
    plt.xlabel('Iteration')
    plt.ylabel('Cost')
    plt.title('Cost Function Convergence')
    plt.grid(True)
    plt.show()
```

Worksheet-5.ipynb

```
# Execute the main function
if __name__ == "__main__":
    main()

...
Initial weights: [0, 0]
Learning rate: 1e-05
Iterations: 1000
Final Weights: [0.28811359 0.64616558]
Cost History (First 10 iterations): [0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
RMSE on Test Set: 27.87654122196794
R-Squared on Test Set: 0.8886354462786421
```

Cost Function Convergence

Cost

Iteration

Worksheet-5.ipynb

```
Iteration
```

Task-11

Performance Assessment:  $R^2 = 0.8$  is very strong for educational data, showing the model explains 80% of variance in Writing scores.

Feature Effectiveness: Both Math and Reading scores are highly predictive of Writing ability.

Learning Rate Suitability: The current learning rate works well given the excellent convergence.

Optimal Model Complexity: Not overfitting (generalizes well) and not underfitting (high explanatory power).

Practical Usefulness: Despite RMSE of 27.8, the high  $R^2$  indicates the model is useful for predictions within the educational context.